

Team 6: LU Factorization

Optimizations targeting towards multicore processors

Ryan McKenna, Matthew Paul, James Kerrigan
Niko Gerassimakis, Neil Duffy

May 20, 2015



- The quintessential problem in linear algebra is solving a linear system of equations

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$$

- We want to find values of x_1 , x_2 , and x_3 such that
 - 1 $a_{11}x_1 + a_{12}x_2 + a_{13}x_3 = b_1$
 - 2 $a_{21}x_1 + a_{22}x_2 + a_{23}x_3 = b_2$
 - 3 $a_{31}x_1 + a_{32}x_2 + a_{33}x_3 = b_3$

- Linear algebra comes up in a lot of professions:
 - Physics
 - Partial differential equations
 - Graph theory
 - Statistics / Curve Fitting
 - Sports Ranking

Solving Linear Systems

- If A is an $n \times n$ matrix, solving a system of the form $Ax = b$ takes $O(n^3)$ time.
- If A is a triangular matrix, then solving the system takes $O(n^2)$ time



LU Factorization Background

- LU Factorization works by decomposing a square matrix A into a lower triangular matrix, L , and an upper triangular matrix, U :

$$A = LU$$


$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} = \begin{bmatrix} l_{11} & 0 & 0 \\ l_{21} & l_{22} & 0 \\ l_{31} & l_{32} & l_{33} \end{bmatrix} \begin{bmatrix} u_{11} & u_{12} & u_{13} \\ 0 & u_{22} & u_{23} \\ 0 & 0 & u_{33} \end{bmatrix}$$

- With L and U , we can solve $Ax = LUx = b$ in $O(n^2)$.



Algorithm Description / Access Pattern

LU factorization is an $O(n^3)$ algorithm:



Implementation

```
void lu(double **A, double **L, double **U, int n) {  
    zero (L, n);  
    copy (U, A, n);  
    init (L, n);  
    for(int j=0; j < n; j++) {  
        for(int i=j+1; i < n; i++) {  
            double m = U[i][j] / U[j][j];  
            L[i][j] = m;  
            for(int k=j; k < n; k++)  
                U[i][k] -= m*U[j][k];  
        }  
    }  
}
```





Approach

- 1 Generate random matrices up to 6400×6400 .
- 2 Run and time 4 trials of the factorization algorithm for each matrix size.
- 3 Repeat for every optimization configuration.



Optimizations

- -O1, -O2, -O3
- loop unrolling
- vectorization
- native (architecture specific) optimizations
- openMP



Compiler



Sequential Results

| | Optimization / Time (s) | | | | | | |
|------|-------------------------|---------|--------|--------|---------|--------|--------|
| n | None | -O1 | -O2 | Vector | -O3 | Unroll | Native |
| 100 | 0.0028 | 0.0010 | 0.0008 | 0.0005 | 0.0005 | 0.0004 | 0.0003 |
| 200 | 0.0101 | 0.0034 | 0.0026 | 0.0017 | 0.0018 | 0.0025 | 0.0014 |
| 400 | 0.07902 | 0.0219 | 0.0189 | 0.0158 | 0.01561 | 0.0151 | 0.0124 |
| 800 | 0.6222 | 0.1575 | 0.1070 | 0.0850 | 0.0821 | 0.0800 | 0.0662 |
| 1600 | 4.9364 | 1.3342 | 0.9388 | 0.7737 | 0.7900 | 0.7465 | 0.6732 |
| 3200 | 39.7114 | 11.7545 | 9.5557 | 8.4604 | 8.4156 | 8.1363 | 8.0514 |
| 6400 | 316.3 | 94.23 | 78.12 | 69.35 | 68.70 | 66.54 | 66.76 |

- Going from no optimizations to -O1 gave nearly 400% speedup.
- Vectorization yielded an 11% speedup.
- Loop unrolling gave a 3 % speedup in the largest case.
- Native optimizations yielded a 15 % speedup in the best case (n = 1600).



Sequential Results

